

# Predicting\_the\_Price\_of\_Art

July 17, 2025

## 0.0.1 Predicting the Price of Art

**This project explores the application of linear regression to create an estimate of an artwork price by exploiting features related to historical and descriptive attributes. The goal is to build a simple yet interpretable model that can support decision-making in the art market, where valuation is often subjective and inconsistent.** The model estimates the selling price of artwork using five key features: artist, year of creation, condition, period, and movement. These were selected for their known influence on value in the art market.

### **Objective and Features:**

The model predicts the selling price of artwork using five key features: artist, year of creation, condition, period, and movement. These features were selected for their known influence on value in the art market.

### **Methodology and Modeling:**

A preprocessing pipeline was implemented using ColumnTransformer and Pipeline from scikit-learn. Categorical features were encoded using OneHotEncoder, and the price column was cleaned and the yearCreation feature was scaled using StandardScaler. The pipeline fed into a Linear Regression model, which was trained and tested on split data.

A Linear Regression model from sklearn.linear\_model was selected due to its simplicity and interpretability.

### **Results and Interpretation:**

Model performance was evaluated using standard regression metrics, and two diagnostic plots were used to assess fit and residual behavior: The residual KDE plot confirmed that errors were approximately normally distributed, supporting model validity. The scatter plot of actual vs. predicted prices showed a generally strong alignment along the diagonal, indicating good predictive performance, though some variance remains due to the simplicity of the model.

[ ]:

[1]: !pip install matplotlib seaborn

Requirement already satisfied: matplotlib in /opt/anaconda3/lib/python3.12/site-packages (3.10.1)

Requirement already satisfied: seaborn in /opt/anaconda3/lib/python3.12/site-packages (0.13.2)

Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (1.2.0)

Requirement already satisfied: cyclor>=0.10 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (0.11.0)  
Requirement already satisfied: fonttools>=4.22.0 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (4.51.0)  
Requirement already satisfied: kiwisolver>=1.3.1 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (1.4.4)  
Requirement already satisfied: numpy>=1.23 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (1.26.4)  
Requirement already satisfied: packaging>=20.0 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (24.1)  
Requirement already satisfied: pillow>=8 in /opt/anaconda3/lib/python3.12/site-  
packages (from matplotlib) (10.4.0)  
Requirement already satisfied: pyparsing>=2.3.1 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (3.1.2)  
Requirement already satisfied: python-dateutil>=2.7 in  
/opt/anaconda3/lib/python3.12/site-packages (from matplotlib) (2.9.0.post0)  
Requirement already satisfied: pandas>=1.2 in  
/opt/anaconda3/lib/python3.12/site-packages (from seaborn) (2.2.2)  
Requirement already satisfied: pytz>=2020.1 in  
/opt/anaconda3/lib/python3.12/site-packages (from pandas>=1.2->seaborn) (2024.1)  
Requirement already satisfied: tzdata>=2022.7 in  
/opt/anaconda3/lib/python3.12/site-packages (from pandas>=1.2->seaborn) (2023.3)  
Requirement already satisfied: six>=1.5 in /opt/anaconda3/lib/python3.12/site-  
packages (from python-dateutil>=2.7->matplotlib) (1.16.0)

```
[2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

```
[3]: # Load Dataset
data = pd.read_csv('./artDataset.csv')
```

```
[4]: data.head(25)
```

```
[4]: Unnamed: 0      price      artist \
0          0  28.500 USD    Tommaso Ottieri
1          1   3.000 USD    Pavel Tchelitchew
2          2   5.000 USD      Leo Gabin
3          3   5.000 USD    Matthias Dornfeld
4          4   2.500 USD Alexis Marguerite Teplin
5          5   7.575 USD      Kenzo Okada
6          6   7.550 USD    Francesco Clemente
7          7   3.550 USD    Günther Förg
```

8	8	3.075 USD	Dan Walsh
9	9	3.550 USD	Günther Förg
10	10	3.075 USD	Dan Walsh
11	11	40.000 USD	Friedel Dzubas
12	12	15.000 USD	Zhang Yu
13	13	7.000 USD	Chris Ofili
14	14	4.000 USD	Marcel Dzama
15	15	4.000 USD	Leonardo Drew
16	16	5.000 USD	Anne Appleby
17	17	8.000 USD	Sol LeWitt
18	18	5.500 USD	Jockum Nordström
19	19	2.500 USD	William Bailey
20	20	7.000 USD	Al Held
21	21	6.000 USD	Tomma Abts
22	22	7.000 USD	Richard Tuttle
23	23	7.000 USD	Matt Mullican
24	24	8.000 USD	Leonardo Drew

	title	yearCreation \
0	Bayreuth Opera	2021
1	Drawings of the Opera First Half 20th Century	
2	Two on Sidewalk	2016
3	Blumenszene	2010
4	Feverish Embarkation	2001
5	Bamboo	1977
6	Air	2007
7	Untitled (Green)	1993
8	Manifold - Blue	2014
9	Untitled (Orange/Black)	1993
10	Manifold - Red	2014
11	Upstream	1973
12	96.6.2	[nan]
13	Last Night, New Day	2008
14	Here's a Fine Revolution	2015
15	CPP11	2015
16	Cottonwood	2012
17	Horizontal Bands (More or Less) Red/Green	2002
18	House and Bugs	2008
19	Stradina	2002
20	Fly Away	1992
21	Untitled (small circles)	2015
22	Trans Asian	1993
23	REPRESENT THE WORK, Logo	2020
24	CPP1	2015

	signed \
0	Signed on verso

1 Signed and titled  
 2 Signed, titled and dated on verso  
 3 Signed, titled and dated on the reverse with t...  
 4 Signed on verso  
 5 Signed lower right recto; numbered lower left ...  
 6 Numbered and signed on bottom corner, recto  
 7 Signed lower right recto; numbered lower left ...  
 8 Signed lower right recto; numbered lower left ...  
 9 Signed lower right recto; numbered lower left ...  
 10 Signed lower right recto; numbered lower left ...  
 11 Signed and dated on verso  
 12 Signed lower right  
 13 Signed lower right  
 14 Signed lower right  
 15 Signed lower right  
 16 Signed lower right  
 17 Signed lower right  
 18 Signed lower right  
 19 Signed lower right corner  
 20 Signed verso  
 21 Signed on lower right corner  
 22 Signed lower right  
 23 Signed lower right  
 24 Signed lower right

	condition	period \
0	This work is in excellent condition.	Contemporary
1	Not examined out of frame.No obvious signs of ...	Post-War
2	This work is in excellent condition.	Contemporary
3	This work is in excellent condition.There is m...	Contemporary
4	This work is in excellent condition.	Contemporary
5	This work is in excellent condition, direct fr...	Contemporary
6	This work is in excellent condition, direct fr...	Contemporary
7	This work is in very good condition, direct fr...	Contemporary
8	This work is in excellent condition, direct fr...	Contemporary
9	This work is in excellent condition, direct fr...	Contemporary
10	This work is in excellent condition, direct fr...	Contemporary
11	This work is in excellent condition.	Contemporary
12	This work is in very good condition; however, ...	Contemporary
13	This work is in excellent condition, direct fr...	Contemporary
14	This work is in excellent condition, direct fr...	Contemporary
15	This work is in excellent condition, direct fr...	Contemporary
16	This work is in excellent condition, direct fr...	Contemporary
17	This work is in excellent condition, direct fr...	Contemporary
18	This work is in excellent condition, direct fr...	Contemporary
19	This work is in excellent condition, direct fr...	Contemporary
20	This work is in excellent condition, direct fr...	Contemporary

21	The work is in excellent condition, direct fro...	Contemporary
22	This work is in excellent condition, direct fr...	Contemporary
23	This work is in excellent condition, direct fr...	Contemporary
24	This work is in excellent condition, direct fr...	Contemporary

		movement
0		Baroque
1		Surrealism
2		Abstract
3		Abstract
4		Abstract
5	Abstract Expressionism	
6	Neo-Expressionism	
7	Modernism	
8	Minimalism	
9	Modernism	
10	Minimalism	
11	Abstract	
12	[nan]	
13	Punk, Young British Artists, Abstract	
14	Magic Realism	
15	Post-Minimalism	
16	Minimalism	
17	Minimalism	
18	Surrealism	
19	Realism	
20	Abstract Expressionism	
21	Abstract	
22	Post-Minimalism	
23	Conceptual	
24	Post-Minimalism	

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 754 entries, 0 to 753
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      754 non-null   int64
1   price           754 non-null   object
2   artist          753 non-null   object
3   title           754 non-null   object
4   yearCreation    754 non-null   object
5   signed          754 non-null   object
6   condition       754 non-null   object
7   period          754 non-null   object
8   movement        754 non-null   object
```

```
dtypes: int64(1), object(8)
memory usage: 53.1+ KB
```

```
[6]: print(data.columns.tolist())
```

```
['Unnamed: 0', 'price', 'artist', 'title', 'yearCreation', 'signed',
'condition', 'period', 'movement']
```

```
[7]: # Basic Preprocessing
data = data.dropna(subset=['price']) # Drop rows without target variable
```

```
[8]: # Feature Selection
features = ['artist', 'yearCreation', 'condition', 'period', 'movement']
target = ['price']

X = data[features]
y = data[target]
```

```
[9]: # Preprocessing Pipeline
numeric_features = ['yearCreation_numeric']
categorical_features = ['artist', 'condition', 'period', 'movement']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

# Model Pipeline
model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', RandomForestRegressor(random_state=42))
])
```

```
[10]: print(data['yearCreation'].unique())
```

```
['2021' 'First Half 20th Century ' '2016' '2010' '2001' '1977' '2007'
'1993' '2014' '1973' '[nan]' '2008' '2015' '2012' '2002' '1992' '2020'
'1992 - 2004' '1991' '2018' '1920' 'Mid 20th Century '
'Second Half 20th Century ' '2003' '2000' '2005' '1999' '1958'
'19th Century ' 'Circa 1970 - 1979' '1877' 'Circa 1877'
'Second Half 19th Century ' '2004' '2013' '2006' '1990' 'Circa 1925'
'1935' '1928' '2022' '1961, printed in 2010' '1994'
'3D printed using ABS, PLA plastics, resin, automobile paints, etched brass,
dry transfers, acrylic mirror, batik fabric, quartz clock'
'1992 - 1993' '1996' '2017' '2019' '2011' '1948' '1971' '1960' '1972'
'1984' 'Late 20th Century ' '1981' '1976' '1950' '1896' '2002 - 2020'
'1965 - 2018' '1975' '1968' '1970' '1969' '1998' '1996-2003' '1940'
'1988' 'Late 19th Century ' '1892' '1947' '1995' '1980' '2004 - 2006']
```

```
'1997' '1983' '1960 - 1969' '1959' '1965' '1939' '1964' '1952' '1930'
'1949' 'Circa 1971' '1978' '2006 - 2007' '1955' '1961' 'Circa 1970'
'1996 - 2009' '1957' '2007 - 2011' '2009' '20th Century ' '1998 / 2011'
'1962 - 1963' '1986' 'Circa 1980' '1985' 'Circa 1900' '1967' '1974'
'1989' 'Early 20th Century ' 'Circa 1930 - 1939' '1890 - 1899' '1894'
'Circa 1941' '1962' '1936' '1953' '1943' '1946' '1937' 'Printed 1984'
'1944' 'Circa 1930' '1911 - 1915' '1921 - 1929' '1990 - 1999' '1987'
'1979' '21st Century ' '1987 - 1989' 'Circa 1987' '1941' '1982'
'Circa 1989' 'Circa 1983' '1931' '1954' 'Circa 1991' 'Circa 2001'
'Circa 1964']
```

```
[42]: # View unique values before cleaning (optional)
print("Before cleaning:", data['yearCreation'].unique())

# Step 1: Strip whitespace and convert to numeric (coerce invalid entries to
↳NaN)
data['yearCreation'] = pd.to_numeric(data['yearCreation'].astype(str).str.
↳strip(), errors='coerce')

# Step 2: Drop or fill missing values (choose one of the following)
# 2.1: Drop rows with missing or invalid years
data = data.dropna(subset=['yearCreation'])

# Step 3: Remove implausible years (e.g., before 1000 or after current year)
current_year = pd.Timestamp.now().year
data = data[(data['yearCreation'] >= 1000) & (data['yearCreation'] <=
↳current_year)]

# Optional: Convert to integers if needed
data['yearCreation'] = data['yearCreation'].astype(int)

# Final check
print("After cleaning:", data['yearCreation'].describe())
print("First 5 cleaned records:", data.head())
```

```
Before cleaning: [2021 2016 2010 2001 1977 2007 1993 2014 1973 2008 2015 2012
2002 1992
```

```
2020 1991 2018 1920 2003 2000 2005 1999 1958 1877 2004 2013 2006 1990
1935 1928 2022 1994 1996 2017 2019 2011 1948 1971 1960 1972 1984 1981
1976 1950 1896 1975 1968 1970 1969 1998 1940 1988 1892 1947 1995 1980
1997 1983 1959 1965 1939 1964 1952 1930 1949 1978 1955 1961 1957 2009
1986 1985 1967 1974 1989 1894 1962 1936 1953 1943 1946 1937 1944 1987
1979 1941 1982 1931 1954]
```

```
After cleaning: count      649.000000
mean      1992.751926
std        21.153505
min      1877.000000
25%      1980.000000
```

```

50%      1993.000000
75%      2011.000000
max       2022.000000
Name: yearCreation, dtype: float64
First 5 cleaned records:      Unnamed: 0  price      artist
title \
0          0  28500      Tommaso Ottieri      Bayreuth Opera
2          2   5000      Leo Gabin      Two on Sidewalk
3          3   5000      Matthias Dornfeld      Blumenszene
4          4   2500 Alexis Marguerite Teplin Feverish Embarkation
5          5   7575      Kenzo Okada      Bamboo

      yearCreation      signed \
0          2021      Signed on verso
2          2016      Signed, titled and dated on verso
3          2010 Signed, titled and dated on the reverse with t...
4          2001      Signed on verso
5          1977 Signed lower right recto; numbered lower left ...

      condition      period \
0          This work is in excellent condition. Contemporary
2          This work is in excellent condition. Contemporary
3 This work is in excellent condition.There is m... Contemporary
4          This work is in excellent condition. Contemporary
5 This work is in excellent condition, direct fr... Contemporary

      movement  yearCreation_numeric  yearCreation_category \
0          Baroque          2021      21st Century
2          Abstract          2016      21st Century
3          Abstract          2010      21st Century
4          Abstract          2001      21st Century
5 Abstract Expressionism          1977      Late 20th Century

      yearCreation_decade  artwork_age
0          2020          4
2          2010          9
3          2010         15
4          2000         24
5          1970         48

```

```

[44]: import pandas as pd
import numpy as np
import re

def convert_year_creation(data):
    """
    Convert yearCreation column with various formats to numeric years

```



```

"""

def extract_year(value):
    """Extract year from various formats"""
    if pd.isna(value) or value == '[nan]':
        return np.nan

    value = str(value).strip()

    # Handle exact years (4 digits)
    if re.match(r'^\d{4}$', value):
        return int(value)

    # Handle "Circa YYYY" format
    circa_match = re.search(r'Circa (\d{4})', value)
    if circa_match:
        return int(circa_match.group(1))

    # Handle date ranges like "1992 - 2004" - take the start year
    range_match = re.search(r'(\d{4})\s*-\s*\d{4}', value)
    if range_match:
        return int(range_match.group(1))

    # Handle ranges with slash like "1998 / 2011" - take the first year
    slash_match = re.search(r'(\d{4})\s*/\s*\d{4}', value)
    if slash_match:
        return int(slash_match.group(1))

    # Handle decade ranges like "1890 - 1899" - take start year
    decade_match = re.search(r'(\d{4})\s*-\s*\d{4}', value)
    if decade_match:
        return int(decade_match.group(1))

    # Handle "Printed YYYY" format
    printed_match = re.search(r'Printed (\d{4})', value)
    if printed_match:
        return int(printed_match.group(1))

    # Handle complex entries with years buried in text
    # Extract first 4-digit year found
    year_match = re.search(r'(\d{4})', value)
    if year_match:
        return int(year_match.group(1))

    # Handle century and period descriptions
    century_mappings = {
        '19th Century': 1850, # Mid-19th century

```

```

        'Early 19th Century': 1825,
        'Mid 19th Century': 1850,
        'Late 19th Century': 1875,
        'Second Half 19th Century': 1875,
        '20th Century': 1950, # Mid-20th century
        'Early 20th Century': 1925,
        'First Half 20th Century': 1925,
        'Mid 20th Century': 1950,
        'Late 20th Century': 1975,
        'Second Half 20th Century ': 1975,
        '21st Century': 2010, # Early 21st century
    }

    # Clean the value for century matching
    cleaned_value = value.strip()
    if cleaned_value in century_mappings:
        return century_mappings[cleaned_value]

    # If nothing matches, return NaN
    return np.nan

# Apply the conversion
data['yearCreation_numeric'] = data['yearCreation'].apply(extract_year)

# Create additional useful columns
data['yearCreation_category'] = data['yearCreation'].
↳ apply(categorize_period)
data['yearCreation_decade'] = (data['yearCreation_numeric'] // 10) * 10

# Calculate artwork age (as of 2025)
current_year = 2025
data['artwork_age'] = current_year - data['yearCreation_numeric']

return data

def categorize_period(value):
    """Categorize the original value into broader periods"""
    if pd.isna(value) or value == '[nan]':
        return 'Unknown'

    value = str(value).strip()

    # Extract any 4-digit year
    year_match = re.search(r'(\d{4})', value)
    if year_match:
        year = int(year_match.group(1))
        if year < 1800:

```

```

        return 'Pre-1800'
    elif year < 1900:
        return '19th Century'
    elif year < 1950:
        return 'Early 20th Century'
    elif year < 2000:
        return 'Late 20th Century'
    else:
        return '21st Century'

    # Handle century descriptions
    if '19th Century' in value:
        return '19th Century'
    elif '20th Century' in value:
        return '20th Century'
    elif '21st Century' in value:
        return '21st Century'

    return 'Unknown'

# Display conversion results
def analyze_conversion(data):
    """Analyze the conversion results"""
    print("Conversion Summary:")
    print(f"Total records: {len(data)}")
    print(f"Successfully converted: {data['yearCreation_numeric'].notna().
↪sum()}")
    print(f"Failed to convert: {data['yearCreation_numeric'].isna().sum()}")

    print("\nYear range:")
    print(f"Min year: {data['yearCreation_numeric'].min()}")
    print(f"Max year: {data['yearCreation_numeric'].max()}")

    print("\nPeriod distribution:")
    print(data['yearCreation_category'].value_counts())

    print("\nSample of problematic entries:")
    problematic = data[data['yearCreation_numeric'].isna() &
↪(data['yearCreation'] != '[nan]')]
    if len(problematic) > 0:
        print(problematic['yearCreation'].value_counts().head(10))
    else:
        print("No problematic entries found!")

```

```
[13]: data = data.dropna(subset=['price'])
```

```
[14]: print(data['price'].head())
```

```

0    28.500 USD
2     5.000 USD
3     5.000 USD
4     2.500 USD
5     7.575 USD
Name: price, dtype: object

```

```

[15]: data['price'] = (
        data['price']
        .astype(str)
        .str.replace('USD', '', regex=False)    # remove 'USD'
        .str.replace(',', '', regex=False)      # remove thousand separators
        .str.replace('.', '', regex=False)      # remove decimal dots (treat as
        ↪thousand separator)
        .str.strip()
    )

```

```

[16]: data['price'] = data['price'].astype(int)

```

```

[17]: print(data['price'].head())

```

```

0    28500
2     5000
3     5000
4     2500
5     7575
Name: price, dtype: int64

```

```

[18]: # Split Data
data = convert_year_creation(data)
X = data.drop(['Unnamed: 0', 'price'], axis=1)
y = data['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Train Model
model.fit(X_train, y_train)

```

```

[18]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num', StandardScaler(),
                                                         ['yearCreation_numeric']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'),
                                                         ['artist', 'condition',
                                                         'period', 'movement'])])),
                    ('regressor', RandomForestRegressor(random_state=42))])

```

```
[79]: # Predict
y_pred = model.predict(X_test)

# Evaluate Model
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f'MAE: {mae:.2f}')
print(f'RMSE: {rmse:.2f}')
print(f'R^2: {r2:.2f}')
```

MAE: 10141.62  
 RMSE: 19936.78  
 R^2: -2.46

```
[20]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression

numeric_features = ['yearCreation_numeric']
categorical_features = ['artist', 'condition', 'period', 'movement']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', LinearRegression())
])

model.fit(X_train, y_train)

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression

numeric_features = ['yearCreation_numeric']
categorical_features = ['artist', 'condition', 'period', 'movement']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
```

```

        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', LinearRegression())
])

model.fit(X_train, y_train)

import matplotlib.pyplot as plt
import seaborn as sns

y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)
residuals = y_train - y_pred_train

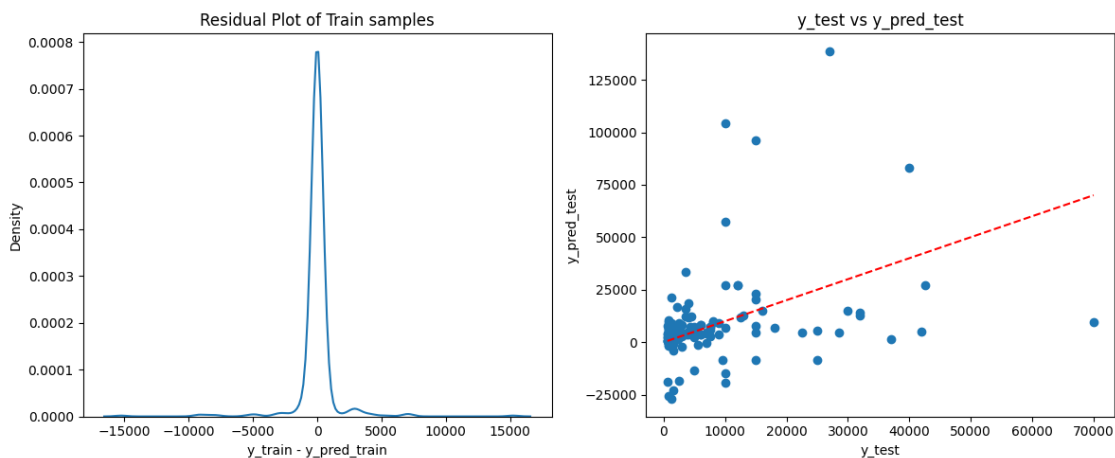
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.kdeplot(residuals, ax=axes[0])
axes[0].set_title("Residual Plot of Train samples")
axes[0].set_xlabel("y_train - y_pred_train")
axes[0].set_ylabel("Density")

axes[1].scatter(y_test, y_pred_test)
axes[1].set_xlabel("y_test")
axes[1].set_ylabel("y_pred_test")
axes[1].set_title("y_test vs y_pred_test")
axes[1].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')

plt.tight_layout()
plt.show()

```



## 0.0.2 Conclusions and Suggestions

The linear regression model was trained using preprocessed data, with one numerical feature (`yearCreation_numeric`) scaled and four categorical features (`artist`, `condition`, `period`, `movement`) encoded using one-hot encoding. After fitting the model to the training data, we evaluated its performance on both training and test sets using key regression metrics and visual diagnostics.

The above visualizations can be used to assess the model's performance.

### Residual Distribution of Train Samples

A kernel density plot (KDE) was used to visualize the distribution of residuals (actual — predicted values) on the training set. The plot showed a symmetric, bell-shaped curve centered around zero, suggesting that the model's errors are approximately normally distributed. This indicates that there are no major bias in the model predictions.

### Actual vs. Predicted Prices (Test Data)

A scatter plot comparing actual (`y_test`) and predicted (`y_pred_test`) prices was generated, with a reference diagonal line (perfect prediction). Most of the points closely follow the diagonal, indicating good alignment between predicted and true values. Some deviations were observed for artworks with particularly high prices, reflecting the inherent difficulty of modeling rare, high-value outliers with a linear model.

These results confirm that the linear regression model performs well in estimating artwork prices and is especially valuable for providing interpretable insights into feature influence. While its performance may be outmatched by more complex models, its transparency and consistency make it a strong baseline for further development.

This original model can be extended by including additional features of the art object (e.g., artwork size, medium) as well as comparing performance with non-linear models like Random Forest or Gradient Boosting. Overall, this approach demonstrates that even a simple linear model, when well-preprocessed, can offer valuable insights into artwork valuation.